Algorithm Design 21/22

Hands On 5 - Bloom Filters

Federico Ramacciotti

1 Problem

1. Consider the Bloom filters where a single random universal hash random function $h : U \to [m]$ is employed for a set $S \subseteq U$ of keys, where U is the universe of keys.

   Consider its binary array B of m bits. Suppose that $m \geq c|S|$, for some constant $c > 1$, and that both c and $|S|$ are unknown to us.

   Estimate the expected number of 1s in B under a uniform choice at random of $h \in \mathcal{H}$. Is this related to $|S|$? Can we use it to estimate $|S|$?

2. Consider B and its rank function: show how to use extra $O(m)$ bits to store a space-efficient data structure that returns, for any given i, the following answer in constant time:

   $rank(i) = \#1s \in B[1..i]$

   Hint: easy to solve in extra $O(m \log m)$ bits. To get $O(m)$ bits, use prefix sums on B, and sample them. Use a lookup table for pieces of B between any two consecutive samples.

2 Solution

2.1 Question 1

Define the indicator variable

$$X_i = \begin{cases} 1 & \text{if } B[i] = 1 \\ 0 & \text{otherwise} \end{cases}$$

The probability for it to be 1 is $Pr[X_i = 1] = Pr[B[i] = 1] = 1 - p' = 1 - \left(1 - \frac{1}{m}\right)^{|S|}$.
Define also the sum $Y = \sum_{i=0}^{m-1} X_i$ and find its expected value:

$$\begin{aligned} E[Y] &= E\left[\sum_{i=0}^{m-1} X_i\right] = \sum_{i=0}^{m-1} Pr[X_i = 1] \\ &= \sum_{i=0}^{m-1} 1 - \left(1 - \frac{1}{m}\right)^n = m - \sum_{i=0}^{m-1}\left(1 - \frac{1}{m}\right)^n \\ &= m - m\left(1 - \frac{1}{m}\right)^n \cong m - me^{-\frac{n}{m}} \\ &= m\left(1 - e^{-\frac{n}{m}}\right) \end{aligned}$$

So we get that the expected number of 1s in B is $m(1-e^{-\frac{n}{m}})$ and, as we can clearly see, it's strongly related to $|S| = n$. In order to estimate $|S|$, given the number of ones $n_1 = E[Y] = m(1-e^{-\frac{n}{m}})$, we get $1-\frac{n_1}{m} = e^{-\frac{n}{m}}$ and, changing the sign and applying the logarithm, we obtain $n = -m \, \log(1-\frac{n_1}{m})$.

2.2 Question 2

2.2.1 Baseline solution

The baseline solution uses prefix sums in $O(m)$ time on the array B to solve any query in costant time. The space used is $O(\log m)$ bits for each integer of the array and, with an array of length m, the total space is $O(m \log m)$.

2.2.2 Better solution: sampling

In order to reduce the space used by the algorithm, we can simply sample the prefix sums array created in the baseline solution. Since the goal is to use linear space $O(m)$ with samples of size $O(\frac{\log m}{2})$, we need just

$$\# \text{ samples} * \frac{\log m}{2} = m \implies \# \text{ samples} = \frac{2m}{\log m}$$

samples. With this solution the query time is no longer constant, because we have to find the right sample in $O(1)$ and then check the interval between two samples in $O(\frac{\log m}{2})$ time.

We can improve this solution to have a constant time per query by pre-computing a lookup table T that stores the prefix sums for any possible integer of $\frac{\log m}{2}$ bits, written in binary. Thus, to get $rank(i)$, we sum the first smaller sample with the right entry in the lookup table T. With $m = 64$ and samples of size $(\log 64)/2 = 3$, the lookup table T is:

$$T = \begin{array}{c|ccc}
(0 \ \ 0 \ \ 0)_2 & 0 & 0 & 0 \\
(0 \ \ 0 \ \ 1)_2 & 0 & 0 & 1 \\
(0 \ \ 1 \ \ 0)_2 & 0 & 1 & 1 \\
(0 \ \ 1 \ \ 1)_2 & 0 & 1 & 2 \\
(1 \ \ 0 \ \ 0)_2 & 1 & 1 & 1 \\
(1 \ \ 0 \ \ 1)_2 & 1 & 1 & 2 \\
(1 \ \ 1 \ \ 0)_2 & 1 & 2 & 2 \\
(1 \ \ 1 \ \ 1)_2 & 1 & 2 & 3
\end{array}$$

Since a string of $\frac{\log m}{2}$ bits can represent up to $2^{\frac{\log m}{2}} = \sqrt{m}$ distinct integers, the lookup table T takes $\log(\frac{\log m}{2}) * \sqrt{m}$ bits. Finally, we get that this solution takes $\sqrt{m}\frac{\log m}{2}\log(\frac{\log m}{2}) = O(m)$ bits, with $O(1)$ time per query.